

CSE 415 Final Exam

Yuchen Wang

TOTAL POINTS

145 / 200

QUESTION 1

A* Search 20 pts

1.1 (a) DFS order of nodes removed from

OPEN 0 / 2

- 0 pts Correct

- 1 pts Partially Correct

✓ - 2 pts Incorrect

1.2 (b) BFS order of nodes removed from

OPEN 0 / 2

- 0 pts Correct

- 1 pts Partially Correct

✓ - 2 pts Incorrect

1.3 (c) Formula, with quantifier, for

admissibility 2 / 4

- 0 pts Correct

- 2 pts Not using quantifier

✓ - 4 pts Incorrect

+ 2 Point adjustment

1.4 (d) Is hB admissible? Justify. 4 / 4

✓ - 0 pts Correct

- 2 pts Partially correct; I'm very strict here. If you do not say if largest disk is not on peg R, it takes at least 1 or 2 moves, you get minus 2; if you say when the largest disk is on peg R, it takes at least 1 move, you get -1 because it could take 0 as in goal state.

- 4 pts Incorrect

1.5 (e) A* order of nodes removed from

OPEN 2 / 4

- 0 pts Correct

✓ - 2 pts Partially Correct

- 4 pts Incorrect

1.6 (f) Nodes expansions of BFS avoided by

A* 0 / 4

- 0 pts Correct

- 2 pts Partially Correct

✓ - 4 pts Incorrect

QUESTION 2

MDPs: Value Iteration 20 pts

2.1 (a) V1(s) for all s. 4 / 4

✓ - 0 pts Correct

- 1 pts Mostly correct [there is a small error]

- 2 pts Partially correct

- 3 pts Mostly incorrect

+ 1 pts shown a correct approach (partially) to answer

- 4 pts No answer

- 4 pts Incorrect answer

2.2 (b) V2(s) for all s. 5 / 8

- 0 pts Correct

- 5 pts Partial credit - more than half of the states are wrong

- 8 pts No answer

- 7 pts Partial credit - very few states are correct

✓ + 1 pts Shown a (partially) correct approach toward the answer

- 8 pts Incorrect answer

- 1 pts Mostly correct (there is a small error)

- 6 pts Partial credit - few states are correct

✓ - 4 pts About half of the states are incorrect

- 3 pts Several states are incorrect

2.3 (c) V3(s) for selected states. 8 / 8

- 0 pts Correct
 - 8 pts no answer
 - ✓ + 1 pts shown some work towards the answer
 - ✓ + 2 pts D (correct approach)
 - ✓ + 2 pts E (correct approach)
 - ✓ + 1 pts A (correct approach)
 - ✓ + 1 pts F (correct approach)
 - ✓ + 1 pts I (correct answer)
 - 0 pts correct
 - 8 pts incorrect answer
 - ✓ - 8 pts reset rubric
- QUESTION 4**
- Joint Prob. Distr. and Bayes Nets 20 pts**
- 5 pts Some setup looks OK but values are wrong.
 - 4 pts Error(s) from episode 2 leading to bad values in episode 3.
 - 1 pts Missing last update, with value for Q(l, Exit) of 87.5
 - + 2 Point adjustment
 - effort

QUESTION 3

Q-Learning 20 pts

3.1 Episode 1. 2 / 6

- 0 pts Correct
- 6 pts Incorrect or blank
- ✓ - 3 pts i1 is OK, but the rest are missing or incorrect.
 - 1 pts Almost all correct.
 - ✓ - 1 pts No justification for answer.
 - 1 pts First 5 transitions are there, but too large.

3.2 Episode 2. 2 / 6

- 0 pts Correct
- ✓ - 2 pts Although your i1 in episode 2 takes advantage of the 50 computed before, your e2 is not doing that.
 - 1 pts First transition (from F towards D) should give 1. Prev. Q value was 0 in this direction from F.
 - 2 pts e2 should be 24.
 - 6 pts Blank or incorrect.
- ✓ - 2 pts First two updates are missing or incorrect.
 - 2 pts i1 should be 75

3.3 Episode 3. 2 / 8

- 0 pts Correct
- 1 pts 1st update should give -1.5 for Q(F, phi5).
- 2 pts 2nd update should give d3 = 10.5
- 2 pts 3rd update should give e3 = 48.5
- ✓ - 8 pts blank or incorrect.

4.1(a) Numbers of free params in the Bayes Net 4 / 4

- ✓ - 0 pts Correct
- 4 pts should be 2, 3, 3. there are 8 in total
- 2 pts there are two free parameters in A
- 2 pts there are only 3 free parameters in B and C
- 2 pts there are 2+3+3=8 in total

4.2 (b) Num. free params in a corresp. full joint. 2 / 2

- ✓ - 0 pts Correct
- 2 pts should be $3 * 2 * 2 - 1 = 11$

4.3 (c) $P(B=b1)$ 2 / 2

- ✓ - 0 pts Correct
- 2 pts should be $0.2 \cdot 0.15 + 0.3 \cdot 0.23 + 0.4 \cdot 0.62$

4.4 (d) $P(C=c1)$ 2 / 2

- ✓ - 0 pts Correct
- 2 pts should be $0.1 \cdot 0.15 + 0.45 \cdot 0.23 + 0.95 \cdot 0.62$

4.5 (e) $P(A=a1|B=B1)$ 4 / 4

- ✓ - 0 pts Correct
- 4 pts should be $(0.2)(0.15)/P(B = b1) = (0.2)(0.15)/(0.2 \cdot 0.15 + 0.3 \cdot 0.23 + 0.4 \cdot 0.62)$ by bayes rule
- 2 pts denominator is not correct
- 2 pts numerator is not correct, should be $0.2 * 0.15$

4.6 (f) $P(C=c_1|B=b_1)$. 0 / 6

- 0 pts Correct
- ✓ - 6 pts should be $P(C=c_1 | B = b_1) = \text{sum_a } P(C = c_1 | A = a) P(A = a | B = b_1)$
- 0 pts B and C are not independent. They are only conditionally independent given A.

- 2 pts Incorrect: It should be: $1 - P_0(s_1) = 1 - 0.5 = 0.5$

- 1 pts Missing numbers! You at-least need to write the values of all the individual terms

QUESTION 5

Markov Models 20 pts

5.1 (a) $P(R | F)$ 4 / 4

- ✓ - 0 pts Correct
- 4 pts Incorrect

5.2 (b) $P(R_{\{t+2\}} | F_{\{t\}})$ 6 / 6

- ✓ - 0 pts Correct
- 2 pts Formula OK except 2 wrong table entries
- 2 pts Symbolic formula OK but numbers are off
- 1 pts Minor arithmetic error
- 3 pts Nontrivial error in formula
- 1 pts Formula OK except 1 wrong table entry
- 6 pts Incorrect or blank

6.2 (b) $P_1(s_1)$ for $t=1$ 5 / 5

- ✓ - 0 pts Correct
- 1 pts Did not normalize the probabilities: $0.3/(0.3 + 0.1) = 0.75$
- 3 pts Incorrect! Value should be(before normalization) $P_0(s_1)P(e_3|s_1) = 0.5 \cdot 0.6 = 0.3$. Then 0.75 after normalization
- 2 pts Took the wrong emission. Should be e_3
- 1 pts Missing numbers! You at-least need to write the values of all the individual terms
- 5 pts Nothing written

6.3 (c) $P_1(s_2)$ for $t=1$ 5 / 5

- ✓ - 0 pts Correct
- 1 pts Did not normalize the probabilities: $0.1/(0.3 + 0.1) = 0.25$
- 3 pts Incorrect! Value should be(before normalization) $P_0(s_2)P(e_3|s_2) = 0.5 \cdot 0.2 = 0.1$. Then 0.25 after normalization
- 2 pts Took the wrong emission. Should be e_3
- 1 pts Missing numbers! You at-least need to write the values of all the individual terms
- 5 pts Nothing written

6.4 (d) $P_2(s_1)$ for $t=2$ 4 / 4

- ✓ - 0 pts Correct
- 0 pts Work looks correct! Given the incorrect values from previous parts
- 1 pts Did not normalize
- 3 pts Incorrect! Value should be(before normalization) $P(e_1|s_1) \cdot (P_1(s_2)P(s_1|s_2) + P_1(s_1)P(s_1|s_1)) = 0.25 \cdot (0.25 \cdot 0.1 + 0.75 \cdot 0.3) = 0.0625$. Then 0.45 after normalization
- 2 pts Took the wrong emission. Should be e_1
- 1 pts Missing numbers! You at-least need to write the values of all the individual terms
- 4 pts Nothing written

QUESTION 6

Hidden Markov Models 20 pts

6.1 (a) $P_0(s_2)$ for $t=0$ 2 / 2

- ✓ - 0 pts Correct

6.5 (d) P2(s2) for t=2 4 / 4

✓ - 0 pts Correct

- 0 pts Work looks correct! Given the incorrect values from previous parts
- 1 pts Did not normalize
- 3 pts Incorrect! Value should be(before normalization) $P(e1|s2) \cdot (P(s1|s2)P(s2|s2) + P(s2|s2)P(s1|s1)) = 0.1 \cdot (0.25 \cdot 0.9 + 0.75 \cdot 0.7) = 0.075$. Then 0.55 after normalization
- 2 pts Took the wrong emission. Should be e1
- 1 pts Missing numbers! You at-least need to write the values of all the individual terms
- 4 pts Nothing written

QUESTION 7

Perceptrons 20 pts

7.1 (a) Weights and threshold for NAND gate

7 / 7

✓ - 0 pts Correct

- 2 pts Only used one negated weight
- 4 pts Missing both negated weights
- 3 pts Incorrect threshold
- 7 pts No answer

7.2 (b) Training (3 cycles) 9 / 9

✓ - 0 pts Correct

- 3 pts Step 1 incorrect
- 3 pts Step 2 incorrect
- 3 pts Step 3 incorrect
- 9 pts No answer
- 2 pts Treated ground truth as part of the data vector

7.3 (c) Find a possible training vector for given outcome 4 / 4

✓ - 0 pts Correct

- 2 pts Incorrect learning rate
- 4 pts Incorrect
- 4 pts No answer
- 1 pts Correct vector off by constant factor

QUESTION 8

Natural Language Processing 20 pts

8.1 (a) Eleven scores 4 / 4

✓ - 0 pts Correct

- 1 pts Almost correct 9 or 10 out of 11
- 2 pts Some right answers: 6, 7, or 8 correct.
- 3 pts Mostly wrong (3, 4, or 5 correct)
- 4 pts Wrong, or fewer than 3 correct.
- 2 pts Easy logs not taken. Examples were shown in solutions to the A6 problem.

8.2 (b) Parse tree 1 3 / 3

✓ - 0 pts Correct

- 2 pts Illegal use of a production.
- 3 pts Incorrect or blank

8.3 (c) Parse tree 2 3 / 3

✓ - 0 pts Correct

- 3 pts Incorrect or blank

8.4 (d) Total score and probability for parse 1 4 / 4

✓ - 0 pts Correct

- 1 pts Score off by 1 or 2
- 2 pts Score off by more than 2 or missing.
- 2 pts Probability missing, inconsistent with score or incorrectly determined from score.

8.5 (e) Total score and probability for parse 2 4 / 4

✓ - 0 pts Correct

- 1 pts Score is 1 or 2 off
- 2 pts Score is more than 2 off, missing, or incorrect in another way.
- 2 pts Probability is inconsistent with score, missing, or incorrectly determined from score.

8.6 (f) More probable parse 2 / 2

✓ - 0 pts Correct

- 2 pts Incorrect or clearly a guess without all the supporting data.

QUESTION 9

Multiple choice 40 pts

9.1 (a) Graph search: main advantage of

BFS vs DFS 5 / 5

✓ - 0 pts Correct

- 5 pts Incorrect

9.2 (b) Expectimax most appropriate if ... 0 /

5

- 0 pts Correct

✓ - 5 pts Incorrect

9.3 (c) A* fails when 0 / 5

- 0 pts Correct

✓ - 5 pts Incorrect

9.4 (d) IDDFS -- what's NOT an advantage 0

/ 5

- 0 pts Correct

✓ - 5 pts Incorrect

9.5 (e) Perceptron training: worst obstacle 5

/ 5

✓ - 0 pts Correct

- 5 pts Incorrect

9.6 (f) One of Asimov's laws of robotics 5 / 5

✓ - 0 pts Correct

- 5 pts Incorrect

9.7 (g) Not part of Kurzweil's "singularity" 0 /

5

- 0 pts Correct

✓ - 5 pts Incorrect

9.8 (h) Obstacles to using Asimov's laws for
self-driving cars 5 / 5

✓ - 0 pts Correct

- 5 pts Incorrect

CSE 415–Winter 2019 — Final Examination

by the Staff of CSE 415, Winter 2019

INSTRUCTIONS: Read these instructions carefully. Write your full name at the top of this cover page. Use the official version of your name that appears in the course registration. Make sure you have all 12 pages. Also write your name on top of every page or it is possible that you would not receive credit for that page. (We will be removing staples from the exams, and if a page of yours gets lost due to not having your name on it, you could miss credit for that page.) Write your UW student number and your UWNNetID where shown below.

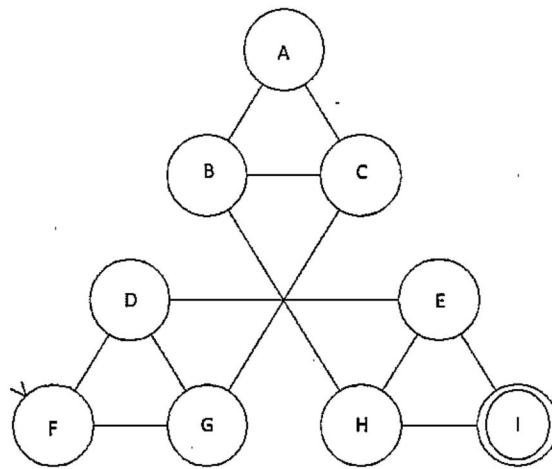
Write legibly and if you use a pencil, make sure that you write darkly enough that a normal scanner will pick up your writing. If you need more space, use the margins. Do all problems. This is a CLOSED-BOOK, CLOSED-NOTES examination. Do not use any books, notes, calculators, or other electronic devices. There are 8 questions worth 20 points each plus one question that has 8 multiple choice parts, and worth 40 points. The total of possible points is 200. Each problem has multiple parts, and the allocations of points among the parts are as shown on each individual problem. Note that some problems may be significantly easier or more time-consuming than others. However, the overall time to complete the exam is estimated to be 90-100 minutes for a student familiar with the material.

POLICY ON QUESTIONS AND CLARIFICATIONS: Do not ask questions during the exam. Instead, if you find one of the exam problems to be ambiguous or unclear, state your objection here on the cover page, clearly identifying the problem number and part number, as well as describing your issue (and continue in the margin on the page of the problem, if necessary). If your objection is accepted by the grader, then points may be adjusted in your score.

UW Student Number:	<u>182672P</u>
UWNNetID	<u>markpb</u>

1 A* Search (Jifan)

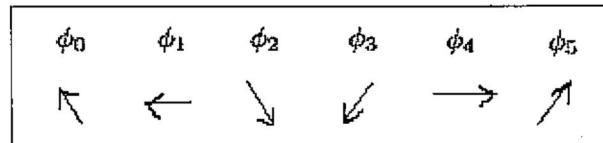
The graph below shows the state space for a 2-disk version of the Towers of Hanoi puzzle. The start state is F (both disks on the left) and the goal state is I. (The top node of the graph is for state A, where both disks are in the middle.) Assume the three pegs of the puzzle are L (left), M (middle), and R (right).



Assume the set of operators is the following, and that they are tried in the given order when generating successors:

- | | |
|---|---|
| 1. ϕ_0 : Try to move a disk from R to M. | 4. ϕ_3 : Try to move a disk from M to L. |
| 2. ϕ_1 : Try to move a disk from R to L. | 5. ϕ_4 : Try to move a disk from L to R. |
| 3. ϕ_2 : Try to move a disk from M to R. | 6. ϕ_5 : Try to move a disk from L to M. |

It might be helpful to think of the operators as moving in particular directions on the graph, according to the following scheme:



- (a) (2 points) What is the order in which nodes are REMOVED from the OPEN list in Depth-First Search?

F D E]

(Just as in the loop-based formulation of DFS presented in lecture, we assume the list L of successors, after deleting nodes already on CLOSED, is pre-pended onto the OPEN list, and any member of L already on OPEN has been removed from OPEN.)

Assume each search (DFS, BFS, A*) terminates as soon as the goal node I is removed from OPEN and found to be a goal node.

- (b) (2 points) What is the order in which nodes are REMOVED from the OPEN list in Breadth-First Search?

F D G E I

- (c) (4 points) In this and the following parts, we'll assume that the cost of each edge in the problem-space graph above is 1. Assume $h(s)$ represents the smallest number of moves needed to get from state s to the goal state (all disks on peg R). Using a formula, give a necessary and sufficient condition for a heuristic function h_i to be admissible. Use a quantifier (either \forall or \exists , as appropriate) in your formula. The set of all states for the puzzle is represented as Σ .

$h_i(\forall s \in \Sigma) \leq \text{actual distance from } s \text{ to goal state.}$

- (d) (4 points) Consider the heuristic function defined as follows:

$$h_B(s) = 0 \text{ if the largest disk is on peg R; 1, otherwise.}$$

Is h_B admissible? Justify your answer.

Yes, it is, because actual distance is 0 if largest disk on R, otherwise

- (e) (4 points) What is the order in which states are removed from the OPEN list by the A* algorithm using heuristic h_B ? (Break ties in priority values by the same ordering from BFS: the order in which states are placed onto the OPEN list.)

F D E I

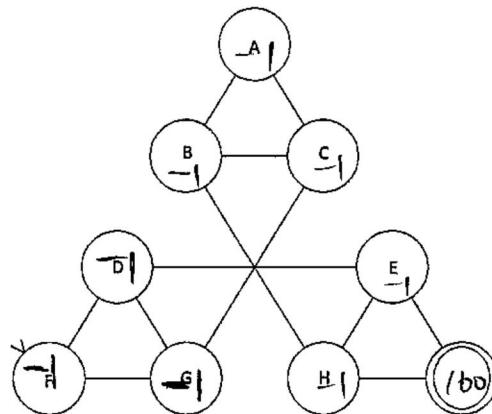
- (f) (4 points) What node expansions, if any, are saved here by A*, versus BFS? (A node expansion is the generation of all that state's successors, whether or not any of those successors already exist.)

2 MDPs: Value Iteration (Hamid)

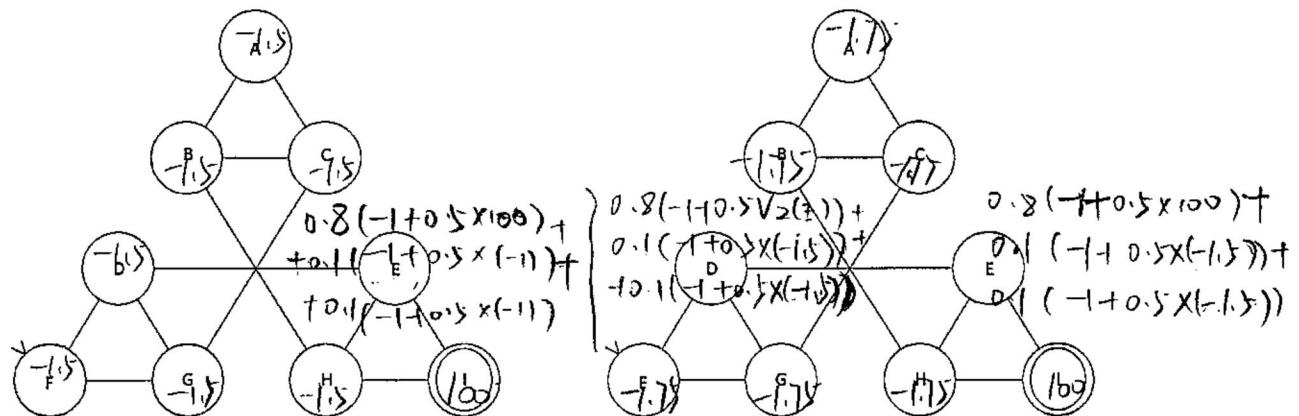
Once again, we are considering the Towers-of-Hanoi puzzle having 2 disks. Our job is now to consider a more general context in which the intended moves sometimes lead to possibly undesired states. Assume we have an MDP whose states are those shown (A through I) plus a terminal state Z (not shown in the diagrams), reached from exit state I. Allowed actions are only those corresponding to graph edges, except that from state I, only the transition to Z is allowed.

All transitions are deterministic except for those starting in D and E, where there is an 0.8 of going in the intended direction and 0.1 probability of going in each of the other two directions. Assume the living cost is $R = -1$ in any transition except I to Z, where the reward is 100. Assume the agent's discount factor is $\gamma = 0.5$. We assume $V_0(s) = 0$ for all s .

- (a) (4 points) Compute $V_1(s)$, and write the values in their states below.



- (b) (8 points) Compute $V_2(s)$, and write the values in their states below on the left copy of the graph.
For $V_2(E)$, an expression involving numbers, multiplications and additions will be sufficient.



- (c) (8 points) Compute $V_3(s)$ for only the following states, and write the values in their states above on the right copy of the graph: A, F, D, E, I. You can just give correct expressions for $V_3(D)$ and $V_3(E)$; we don't require that you calculate the final number.

3 MDPs: Q-Learning (Jifan)

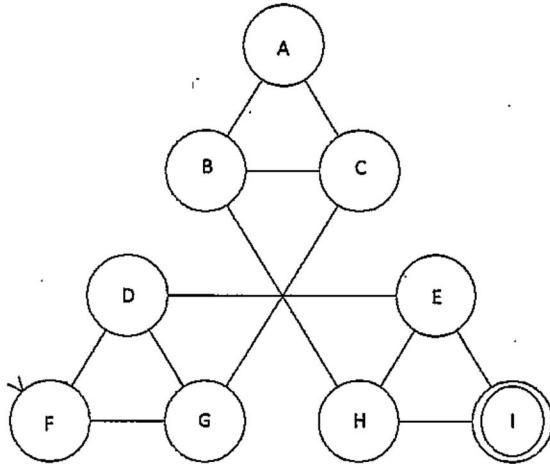
Let's now consider what an agent does in the same Towers-of-Hanoi world (expressed as a Markov Decision Process again), when it is not given the MDP parameters T and R, but must learn about the MDP through its own experience. The agent's discount factor here is: $\gamma = 1$. (No discounting)

Use the following Q-update rule after each transition of the following episodes to adjust the Q values (all initially 0), and show updated values next to the outgoing edges from each node. Use a learning rate $\alpha = 0.5$. For example, the first transition in episode 1 should result in a value of -1 coming out of state F towards G, since the sample has value -2 (because the reward is -2 , the maximum Q value out of G is 0), the existing $Q(F, \phi_4)$ is 0 , and the learning rate is 0.5 .

$$Q(s, a) \leftarrow \alpha \text{ sample} + (1 - \alpha)Q(s, a)$$

$$\text{where sample} = r + \gamma \max_{a'} Q(s', a')$$

Here the transition being processed is $[s, a, s', r]$.



Here are the episodes. Show all updates, but extra space is provided for details on some.

- (6 points) $[F, \phi_4, G, -2]$, $[G, \phi_5, C, -2]$, $[C, \phi_1, B, -2]$, $[B, \phi_2, H, -2]$, $[H, \phi_4, I, -2]$, $[I, \text{Exit}, Z, 100]$

Show your computation for $i_1 = Q(I, \text{Exit})$ after the transition from I.

$$i_1 = \underline{50}$$

- (6 points) $[F, \phi_5, D, -2]$, $[D, \phi_4, E, -2]$, $[E, \phi_2, I, -2]$, $[I, \text{Exit}, Z, 100]$

Show your computation for $e_2 = Q(E, \phi_2)$ after the transition from E.

$$e_2 = \underline{-1}$$

Show your computation for $i_2 = Q(I, \text{Exit})$ after this episode's transition from I.

$$i_1 = 0.5 Q(I, \text{Exit}) + 0.5 (100 + 1 \times 0) = 75$$

- (8 points) $[F, \phi_5, D, -2]$, $[D, \phi_4, \text{Exit}, -2]$, $[E, \phi_2, I, -2]$, $[I, \text{Exit}, Z, 100]$

Show your computation for $d_3 = Q(D, \phi_4)$ after this episode's transition from D.

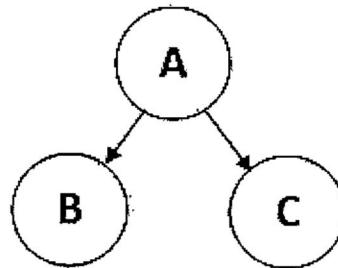
$$d_3 = 0.5 Q(\text{Exit}, \phi_4) + 0.5 (-2 + 1 \times (-1)) = -1.5$$

Show your computation for $e_3 = Q(E, \phi_2)$ after this episode's transition from E.

$$e_3 = 0.5 \times (-1.5) + 0.5 (-2 + 1 \times (-1)) = -2.25$$

4 Joint Probability Distributions and Bayes Nets (Sam)

In this example, we consider a joint probability distribution of three random variables with different domains. We use a Bayes net to represent this distribution. The graph and conditional probability tables for the Bayes net are shown below.



A	P(A)
a ₁	0.15
a ₂	0.23
a ₃	0.62

B	A	P(B A)
b ₁	a ₁	0.2
b ₂	a ₁	0.8
b ₁	a ₂	0.3
b ₂	a ₂	0.7
b ₁	a ₃	0.4
b ₂	a ₃	0.6

C	A	P(C A)
c ₁	a ₁	0.1
c ₂	a ₁	0.9
c ₁	a ₂	0.45
c ₂	a ₂	0.55
c ₁	a ₃	0.95
c ₂	a ₃	0.05

In the following, wherever “compute” is requested, give a formula that shows what numbers and operations are required, but don’t actually do the arithmetic. For example:

$$P(A = a_1, B = b_1 | C = c_1) = (0.15)(0.2)(0.1) / (0.1 + 0.5 + 0.9)$$

(a) (4 points) What are the numbers of free parameters in each of the three tables? 2 3 3
What is the total number of free parameters among the tables? 8

(b) (2 points) If the joint distribution for A, B, and C were general and could not be factored as this Bayes net, how many free parameters would it have? 11

(c) (2 points) “Compute” $P(B = b_1)$. $0.2 \times 0.15 + 0.3 \times 0.23 + 0.4 \times 0.62$

(d) (2 points) “Compute” $P(C = c_1)$. $0.1 \times 0.15 + 0.45 \times 0.23 + 0.95 \times 0.62$

(e) (4 points) “Compute” $P(A = a_1 | B = b_1)$. $\frac{0.2 \times 0.15}{0.2 \times 0.15 + 0.3 \times 0.23 + 0.4 \times 0.62}$

(f) (6 points) “Compute” $P(C = c_1 | B = b_1)$. For this example, you may wish to define one or more other probabilities (and “compute” them) to simplify your main expression.

$$P(C = c_1 | B = b_1) = \frac{P(B = b_1 | C = c_1) P(C = c_1)}{P(B = b_1)} = \frac{P(B = b_1, C = c_1)}{P(B = b_1)}$$

$$P(B = b_1 | C = c_1) = 0.15 \times 0.2 \times 0.1 + 0.23 \times 0.3 \times 0.45 + 0.62 \times 0.4 \times 0.95$$

$$P(C = c_1 | B = b_1) = \frac{0.15 \times 0.2 \times 0.1 + 0.23 \times 0.3 \times 0.45 + 0.62 \times 0.4 \times 0.95}{0.2 \times 0.15 + 0.3 \times 0.23 + 0.4 \times 0.62}$$

5 Markov Models (Sam)

Assume we model the stock market with a two-state Markov model. On each day, the market is either rising (R) or falling (F). The transition matrix is this:

X_t	X_{t+1}	$P(X_{t+1} X_t)$
R	R	0.8
R	F	0.2
F	R	0.3
F	F	0.7

- (a) (4 points) Assuming the market is falling today, determine the probability that it will be rising tomorrow.

$$P(X_{t+1}=R|X_t=F) = 0.3$$

- (b) (6 points) Assuming the market is falling today, determine the probability that it will be rising the day after tomorrow. (It may or may not be rising tomorrow.)

$$P = 0.3 \times 0.8 + 0.7 \times 0.3$$

$$= 0.45$$

- (c) (10 points) Find the stationary probability $r = P_\infty(R)$ that the market is rising.

$$r = P(R|R)r + P(R|F)(1-r)$$

$$r = 0.8 \times r + 0.3(1-r)$$

$$r = 0.6$$

6 Hidden Markov Models (Divye)

In this exercise, we will run the Forward algorithm for 2 steps, for the following given Hidden Markov Model.

Transition model:

X_t	X_{t+1}	$P(X_{t+1} X_t)$
s_1	(s_1)	0.3
s_1	s_2	0.7
s_2	(s_1)	0.1
s_2	s_2	0.9

Emission model:

X_t	e_1	e_2	e_3
s_1	0.25	0.15	0.6
s_2	0.1	0.7	0.2

Let's say the initial probability for s_1 is $P_0(s_1) = 0.5$ (at $t = 0$), and we observed e_3 at $t = 1$ and e_1 at $t = 2$.

In each of the following, an unevaluated expression involving operations and numbers will be sufficient; you don't need a calculator.

(a) (2 points) What is $P_0(s_2)$ for $t = 0$?

$$P_0(s_2) = 0.5$$

(b) (5 points) What is $P_1(s_1)$ (for $t = 1$)? (This is after the observation of e_3 .)

$$P_1(s_1) = P(e_3|s_1) P_0(s_1) = 0.3$$

$$\text{normalize: } P_1(s_1) = 0.75$$

(c) (5 points) Similarly, what is $P_1(s_2)$ (for $t = 1$)?

$$P_1(s_2) = P(e_3|s_2) P_0(s_2) = 0.1$$

$$\text{normalize: } P_1(s_2) = 0.25$$

(d) (4 points) What is $P_2(s_1)$ (for $t = 2$)?

$$P_2(s_1) = P(e_1|s_1)(P(s_1|s_2)P_1(s_2) + P(s_1|s_1)P_1(s_1))$$

$$= 0.25 \times (0.1 \times 0.25 + 0.3 \times 0.75)$$

normalize: $P_2(s_1) = \frac{1}{11}$

(e) (4 points) What is $P_2(s_2)$ (for $t = 2$)?

$$P_2(s_2) = P(e_1|s_2)(P(s_2|s_1)P_1(s_1) + P(s_2|s_2)P_1(s_2))$$

$$= 0.1 (0.7 \times 0.75 + 0.1 \times 0.25)$$

$$\text{normalize: } P_2(s_2) = \frac{6}{11}$$

7 Perceptrons (Rob)

For the following questions, assume all perceptrons output 1 iff their weighted input a satisfies $a > \theta$, and 0 otherwise. (This is not $a \geq \theta$.)

- (a) (7 points) Write the weights and threshold for a two-input perceptron that would cause it to function as a NAND gate (that is, the negation of the logical AND). Possible input values are $\{0, 1\}$, where 0 represents False, and 1 represents True.



- (b) (9 points) Assuming an initial weight vector of $\langle 1, 1 \rangle$ and static $\theta = 0$, train this perceptron with the following example data. Use a learning rate of $\frac{1}{2}$. Write the weight vector that would result after each training example. There is no bias term for this perceptron. Expected output for each example is listed after the data vector:

$$(i) (\langle 0, 0 \rangle, 0) \quad \langle 1, 1 \rangle$$

$$(ii) (\langle -2, 3 \rangle, 1) \quad \langle 1, 1 \rangle$$

$$(iii) (\langle 0, 1 \rangle, 0) \quad \langle 1, 0 \rangle$$

- (c) (4 points) Assume we initially had a perceptron with weight vector $\langle 0, 0, 0 \rangle$ with $\theta = 1$. Then we gave a positive training example that was misclassified as negative and resulted in our weight vector changing to $\langle 1, 0, 2 \rangle$. Give a possible training data vector and learning rate that could cause this result.

Training data:
 $(\langle 1, 0, 2 \rangle, 1)$

Learning rate is $\frac{1}{2}$

8 NLP (Steve)

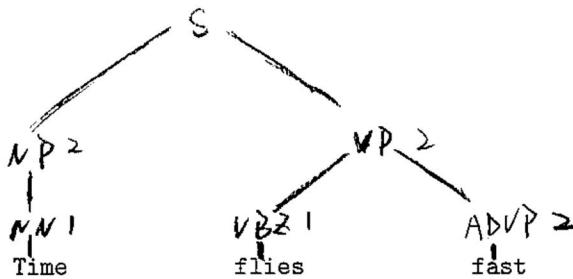
Consider the sentence "Time flies fast." This might mean that the clock seems to move quickly, or it might mean that some kind of insects are not eating.

With the given probabilistic context-free grammar, find two legal parses, and compute a score for each one. Then identify the most probable parse using the scores. The grammar is given below. Consider the number at the right of a production to be the conditional probability of applying that production given that the symbol to be expanded is the symbol on the left-hand side of the production. Convert each probability into a score by taking score = $-\log_2(p)$.

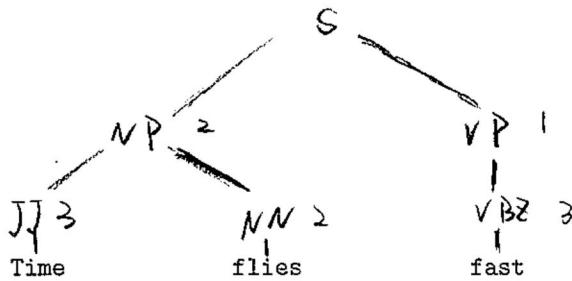
S	$::=$	NP VP	1.00	0
NP	$::=$	JJ NN	0.25	2
NP	$::=$	NN	0.25	2
VP	$::=$	VBZ	0.50	1
VP	$::=$	VBZ ADVP	0.25	2
NN	$::=$	Time	0.50	1
JJ	$::=$	Time	0.125	3
NN	$::=$	flies	0.25	2
VBZ	$::=$	flies	0.50	1
VBZ	$::=$	fast	0.125	3
ADVP	$::=$	fast	0.25	2

(a) (4 point) Scores for each production rule: write the scores next to the probabilities above.

(b) (3 points) Parse number 1:



(c) (3 points) Parse number 2:



(d) (4 points) Total score 8 and overall probability 2^{-8} for parse number 1.

(e) (4 points) Total score 11 and overall probability 2^{-11} for parse number 2.

(f) (2 point) Which parse is more probable? 1

9 Multiple Choice (Rob)

This question is worth 40 points (5 for each subquestion). For each subquestion, underline the letter of the *best* answer. For example: A.

- (a) In graph search with a constant branching factor $b \geq 2$, what is the main advantage of BFS over DFS?
- (A) Guarantee of using less memory to find a shortest path to the goal than with Depth-First Search?
 - (B) Guarantee of identifying a shortest path upon arrival at a goal state.
 - (C) By processing the OPEN list like a queue rather than a stack, the generation of successors naturally goes faster.
 - (D) By searching in all directions equally, there is no chance that the search algorithm will fail to reach a goal if one is reachable.
- (b) Expectimax search is particularly appropriate in adversarial search if
- (A) the opponent is rational but always trying to minimize the state score.
 - (B) the opponent is irrational and unpredictable.
 - (C) the opponent is rational and always trying to maximize the state score.
 - (D) the opponent always makes a predictable move.
- (c) A* search will fail to find a goal node
- (A) if its heuristic is not admissible.
 - (B) only if no goal node is reachable.
 - (C) if its heuristic is not consistent.
 - (D) only if no goal node exists.
- (d) In tree search, which is not an advantage of iterative deepening search over straight depth-first search to the same maximum depth?
- (A) In game-playing, it can provide an anytime algorithm for computing a good move within the available time.
 - (B) When searching for a goal, when it does find a goal, it can immediately identify a shortest path.
 - (C) Its memory requirements are no greater, and on average, less.
 - (D) It might take longer to find a goal node.

- (e) Perceptron training can be viewed as a kind of search for a certain type of hyperplane. What is most likely to prevent finding a solution?
- (A) A learning rate that is constant and equal to 1.
(B) A training set in which the positive and negative examples cannot be completely separated by any hyperplane.
(C) An iterative algorithm that has no limit on the number of iterations it might use during training.
(D) The existence of extra dimensions in the vector representation of the data items, such that these dimensions contain irrelevant features.
- (f) Which of these is one of Asimov's Three Laws of Robotics?
- (A) People must respect the autonomy of any real robot.
(B) A robot must protect its own existence unless that conflicts with the other laws.
(C) People and robots must stay within their own separate realms of expertise.
(D) No robot may consume a resource such as food, which is a resource needed by humans for survival.
- (g) Which of these is not part of Kurzweil's "singularity"?
- (A) Robots are expected to fully control humans.
(B) Robots will have been improving their own designs.
(C) Great strides will have been made in neuroscience.
(D) Artificial intelligence will help spread human values.
- (h) Which of these stand in the way of using Asimov's rules for self-driving cars?
- (A) A car cannot be a robot, and vice versa.
(B) Robots cannot always predict the consequences of their actions accurately.
(C) The rules are written in English and therefore irrelevant to robots and computers.
(D) The rules need to be adopted by state legislatures before they can be applied.